

# QL vs. SARSA: Mathematical Analysis

Joseph Parayil

March 2025

## Claim

Let the training configuration  $C$  be defined as:

$$C = (\Omega, T, E, S, \alpha, \gamma, \epsilon)$$

where:

$\Omega$  is the environment.

$T$  is the number of trials.

$E$  is the number of episodes.

$S$  is the maximum number of steps per episode.

$\alpha$  is the learning rate.

$\gamma$  is the discount factor.

$\epsilon$  is the exploration probability at any given step.

Claim: There exists a training configuration  $C$  such that:

$$P_{QL}(C) > P_{SARSA}(C)$$

There exists a configuration where Q-learning's average convergence performance exceeds SARSA's.

## Assumptions

CAUSAL CHAIN MODEL: The difference between the two algorithm's update rules snowballs into stark differences in performance through the following causal chain:

1. A difference in the update rule causes a difference in the Q-values
2. The maximum Q-value action for each state determines the policy  $\pi$
3. The policy  $\pi$  determines which actions are taken by the agent and which states it then experiences
4. The current policy  $\pi$  and experienced states further affect how the Q-values are updated
5. Steps #1-4 repeat in a feedback loop for a set number of steps and episodes
6. The final average convergence values are averaged over T trials for average convergence score P.

## Proof

### Lemma #1:

SARSA's lower convergence is due to not reaching the goal state sufficiently.

From Gridworld #31:

$$P_{SARSA}(C) = 0.2$$

Since  $P_{SARSA}(C) > 0$ ,

$$\exists i \in [0, n] [P_i > 0]$$

- Since the average performance of SARSA is greater than 0, there are at least some SARSA agents which manage to reach the goal state

Postulate: Performance results  $P_i$  from 0.7 to 1 inclusive are 'goal-seekers' and have found the path to the goal state. Those that are zero or below are 'corner-campers' who have not found the path to the goal state.

$$|P_{SARSA}(C) - 0| < |P_{SARSA}(C) - 0.7|$$

- The average performance of SARSA is closer to 0 than it is 0.7. Therefore, the majority of SARSA agents are 'corner-campers'.

Therefore, although some SARSA agents do manage to find the goal state, the majority of them are not able to. SARSA's lower convergence performance at gridworld #31 is due to not reaching the goal state sufficiently.

## Lemma #2:

For every timestep, SARSA's update rule is exactly equal to QL's update rule with probability  $1 - \epsilon$

First, I will prove that, when  $\epsilon = 0$ , SARSA's update rule is exactly identical to QL's.

### QL update rule:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$$

### SARSA update rule:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

In SARSA, when  $\epsilon = 0$ ,

$$a_{t+1} = \text{argmax}(Q(s_t, a))$$

-The next action is always the maximum q-value action when there is no exploration probability. Substituting this value into the argument of the Q function:

$$Q(s_t, a_t) = Q(s_t, \text{argmax}(Q(s_t, a)))$$

-Based on the definition of argmax (returning index of maximum value) and max (returning maximum value), this can be simplified as follows:

$$Q(s_t, \text{argmax}(Q(s_t, a))) = \max_a Q(s_{t+1}, a)$$

-Therefore:

$$Q(s_t, a_t) = \max_a Q(s_{t+1}, a)$$

- Substituting this into the SARSA expression gives:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$$

which is identical to QL's expression. Therefore, when  $\epsilon = 0$ , SARSA's update rule equals QL's. This can then be extended to say that SARSA's update rule equals QL's with probability  $1 - \epsilon$

### Corollary:

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{n=0}^{\infty} Q(s_t, a_t) = \mathbb{E} \left[ \sum_a \pi(S_{t+1}, a) Q(s_{t+1}, a) \right]$$

-When  $\epsilon = 0.1$ , SARSA's bootstrap expression, on average, approaches the maximum q-value, discounted by 0.925, plus the 3 other possible actions each discounted by 0.25.

### Lemma #3:

Insufficient exploration is the cause of SARSA not reaching the goal state.

Consider Gridworld #31:




Figure 1: Gridworld #31

Feeding the SARSA's expected expression

$$\sum_a \pi(S_{t+1}, a) Q(s_{t+1}, a)$$

into the Value-Iteration algorithm gives us the following optimal policy:




Figure 2: Gridworld #31's optimal policy using SARSA's expected bootstrap expression

-A direct path from start to goal state is observed. -However, as stated in Lemma #1, the majority of SARSA agents did not obtain this policy.

-In Temporal-Differencing algorithms, convergence to the optimal\* policy is guaranteed if sufficient exploration is done by the agent.

\*optimal in this case means treating  $\epsilon$  probability exploration as a transition probability part of the environment itself

Let  $a$  be a SARSA agent of one of the training runs for #31

Let  $E(a)$  mean that sufficient exploration was conducted during the training run

Let  $C(a)$  mean that the agent converged to the optimal policy

Let  $G(a)$  mean that the agent is a goal-seeker and finds the goal

1.  $E(a) \implies C(a)$  [Premise]
2.  $C(a) \implies G(a)$  [Premise]
3.  $\neg G(a)$  [Premise]
4.  $E(a) \implies G(a)$  [Hypothetical Syllogism on (1) and (2)]
5.  $\neg G(a) \implies \neg E(a)$  [Contraposition on (4)]
6.  $\neg E(a)$  [Modus Ponens on (3) and (5)]

- Therefore, SARSA does not explore the environment sufficiently enough in Gridworld #31, while QL does. Causal event #1 is not the sole reason for SARSA's lower convergence. It has to do with Causal event #4 as well.

#### **Lemma #4:**

Information about reward values propagate slowly through the gridworld. The speed at which they propagate is a function of their distance to the start state.

Every episode, an agent starts in the same state  $s_i$ , and follows an epsilon-greedy strategy for moving around a maze. This walking behavior can be loosely modeled as a random-walk.

Due to the random-walk principle, state visits to other states diminish in frequency based on their distance from the start.

Additionally, due to the nature of Temporal-Difference bootstrapping, the Q-values of  $(s, a)$  pairs update based on neighboring  $(s, a)$  pairs only when the agent moves through those states.

Let  $f(d)$  be the average frequency of visits to a state  $d$  spaces away from the start state. If we model the agent's movement as a 2D random-walk, we know that the probability distribution of the different states forms a bell-curve.  $p(d)$ , the average period between each visit to the state, would then be the inverse:

$$p(d) = \frac{1}{f(d)}$$

For the sake of argument, all we have to know is that, as  $d$  gets higher, the period increases dramatically.

The total number of timesteps needed for q-value information on a state  $d$

spaces away to travel to the start state would be

$$t(d) = \sum_{n=1}^d [p(n)]$$

Therefore, q-values closer to the start state get propagated and accounted for much sooner than q-values farther away can.

### **Lemma #5:**

SARSA's evaluation of going to the 'bottleneck' state in Gridworld 31 is a negative Q-value, whereas for QL it is a 0 Q-value. SARSA avoids going to the bottleneck again.

Consider Gridworld #31 again, specifically at state known as  $s_{27}$ :




Figure 3: State  $s_{27}$  in Gridworld #31

Let's imagine a SARSA agent that is encountering this particular state for the first time. The maximum q-value action at that state is the rightward direction. The policy  $\pi$  would be:

$$\begin{aligned}\pi(\uparrow, s_{27}) &= 2.5\% \\ \pi(\downarrow, s_{27}) &= 2.5\% \\ \pi(\leftarrow, s_{27}) &= 2.5\% \\ \pi(\rightarrow, s_{27}) &= 92.5\%\end{aligned}$$

Now imagine the same agent in the next episode, this time in the state  $s_{26}$  to the left:

$$Q(\rightarrow, s_{26}) \rightarrow \gamma \sum_a \pi(S_{27}, a) Q(s_{27}, a)$$

$$Q(\rightarrow, s_{26}) \rightarrow -0.025$$

and

$$Q(\leftarrow, s_{26}) \rightarrow \gamma \sum_a \pi(S_{25}, a) Q(s_{25}, a)$$

$$Q(\leftarrow, s_{26}) \rightarrow 0$$

Since  $Q(\leftarrow, s_{26}) > Q(\rightarrow, s_{26})$ , SARSA's chance of choosing right again at that state would be just 2.5%.

With QL, both evaluations would be zero due to its use of the max function, and QL would be more able to explore the rest of the map by taking a right.

SARSA, on the other hand, stays in the left size area of the map, and does not ever discover the goal state, unless it gets several lucky explorations in a row many different times during training, which is unlikely.

It seems that QL's ignorance of the risks of epsilon-probability exploration made it explore more, leading it to gain more knowledge about the environment than SARSA ever could

QL can beat SARSA because SARSA's accounting of  $\epsilon$ -probability exploration's costs prevents further exploration of the environment, causing it to converge at a suboptimal policy.  $\square$